

MODULE 16

Exception Handling in Distributed APIs

Build resilient APIs by handling errors gracefully across distributed systems.







WHY API ERROR HANDLING MATTERS

Poor error handling leaks internal details, confuses consumers, and makes debugging painful.

POOR ERROR HANDLING

- Inconsistent error formats across endpoints
- Stack traces and internal details exposed
- Generic "500 error" with no useful information
- Hard for consumers to handle programmatically

VS

GOOD ERROR HANDLING

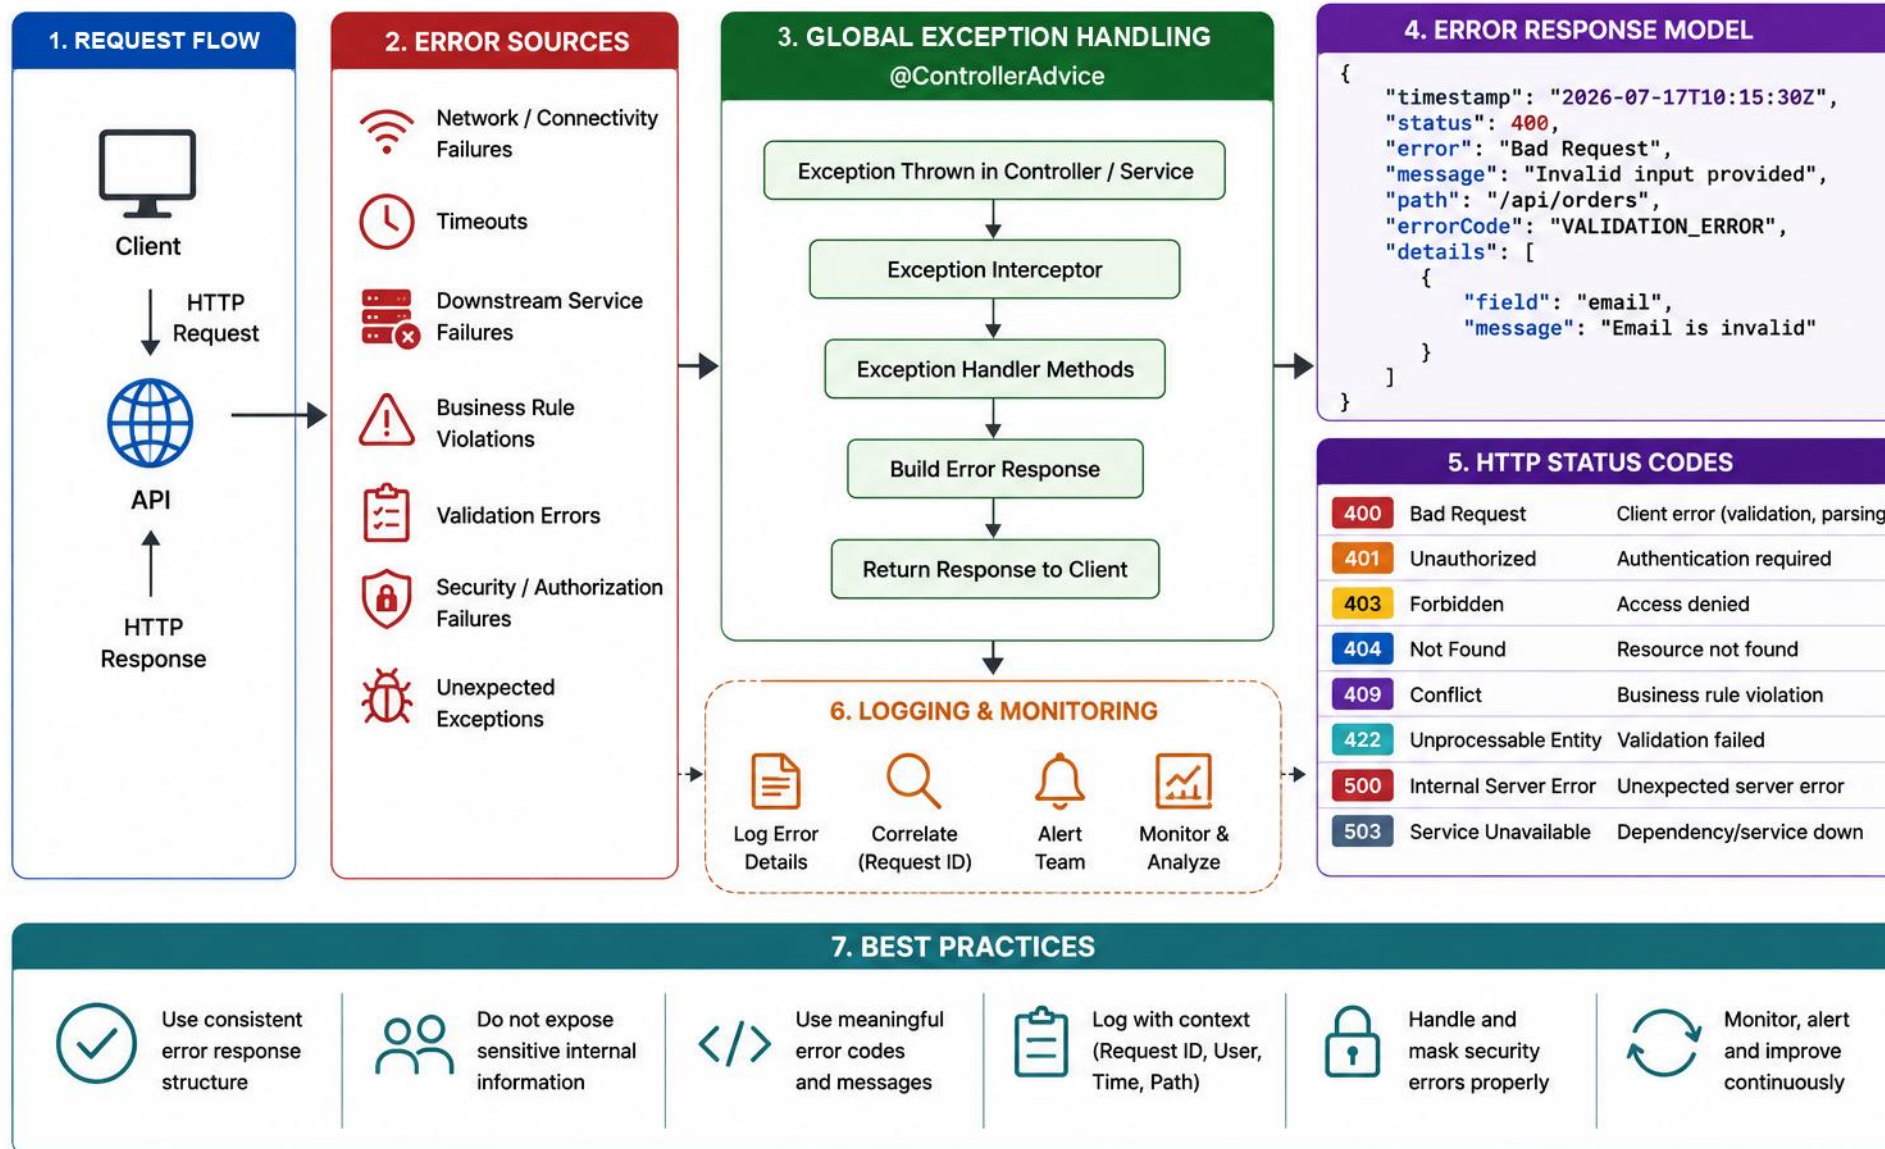
- Consistent, structured error responses
- No sensitive internal details exposed
- Meaningful codes, messages, and correlation IDs
- Easy for consumers to parse and act on

MORE REASONS IT MATTERS

- Improves reliability — graceful handling prevents crashes and cascading failures.
- Enhances user experience — clear, consistent errors show clients what to fix.
- Simplifies client integration and speeds up troubleshooting.
- Builds trust and professionalism: "Errors are not failures — they are opportunities to communicate, guide, and recover."

KEY TAKEAWAY Good error handling is part of your API contract — it builds trust and makes integration easier for consumers.

API ERROR HANDLING



DISTRIBUTED SYSTEM FAILURES

In distributed systems, failures are the norm, not the exception — design for them.



Network Failures

Timeouts, dropped connections, DNS failures.



Latency & Timeouts

Slow downstream services blocking requests.



Service Unavailability

A dependency is down or overloaded.



Partial Failures

Some operations succeed, others fail.



Retries & Duplicates

Retries can cause duplicate side effects.



Cascading Failures


One failure triggers failures elsewhere.

Exception handling communicates failure. Resilience patterns (retry, timeout, circuit breaker, fallback) attempt recovery or containment — covered in detail later.

DOWNSTREAM SITUATION	API BEHAVIOR
Connection failure	502 or 503, depending on context
Timeout	504 when acting as gateway/proxy
Circuit open	Usually 503
Rate limited downstream	Propagate or map 429 carefully
Partial success	Domain-specific result, compensation, or async status
Unknown outcome after timeout	Do not blindly retry non-idempotent operations

KEY TAKEAWAY Design for failure: timeouts, retries with backoff, circuit breakers, and idempotency are essential in distributed APIs.

DISTRIBUTED SYSTEM FAILURES



ERROR CLASSIFICATION

Classify errors by type so each can be handled, logged, and reported appropriately.

CLASS	MEANING	TYPICAL STATUS
Client Input	Malformed, invalid, or missing data	400
Domain Conflict	Business rule or state violation	409
Authentication / Authorization	Identity or permission failure	401 / 403
Resource Lookup	Requested resource absent	404
Dependency Failure	Timeout, unavailable downstream, bad gateway	502 / 503 / 504
Internal Failure	Unexpected application defect	500


CLASSIFICATION DIMENSIONS

- Source: client, domain, dependency, internal. Retryability: retryable, non-retryable, unknown. Exposure: consumer-safe detail vs. internal-only detail.

ERROR	SOURCE	RETRYABLE?	CLIENT DETAIL
Invalid email	Client	No	Field-level detail
Duplicate customer	Domain	No	Stable conflict code
Downstream timeout	Dependency	Maybe	Generic temporary-failure message
Null pointer	Internal	No	Generic server error


KEY TAKEAWAY A clear error classification scheme is the foundation for consistent handling, logging, and HTTP status mapping.

ERROR CLASSIFICATION



GLOBAL REST EXCEPTION HANDLING WITH @RestControllerAdvice

Centralize exception handling so every controller returns consistent, well-formed error responses.



BENEFITS



Cleaner Controllers

No repeated try/catch blocks.



Consistent Responses

Same format everywhere.



Easier Maintenance

One place to update error logic.

```
@ExceptionHandler(CustomerNotFoundException.class) ...  
@ExceptionHandler(MethodArgumentNotValidException.class) ...  
@ExceptionHandler(Exception.class) // fallback, last
```


HANDLER ORDER & SCOPE

Specific handlers first, framework rules next, generic fallback last. Never leak raw exception text; log unexpected failures with a trace/correlation ID.

@ControllerAdvice centralizes cross-controller advice; @RestControllerAdvice combines it with response-body serialization.

KEY TAKEAWAY Global exception handling removes duplicated error-handling code and guarantees consistency across every endpoint.

GLOBAL EXCEPTION HANDLING



HANDLED EXCEPTIONS (EXAMPLES)

ValidationException  @ExceptionHandler 400 Bad Request	ResourceNotFoundException  @ExceptionHandler 404 Not Found	UnauthorizedException  @ExceptionHandler 401 Unauthorized	BusinessException  @ExceptionHandler 422 Unprocessable Entity	ExternalServiceException  @ExceptionHandler 503 Service Unavailable	Exception (Generic)  @ExceptionHandler 500 Internal Server Error
---	---	--	--	--	---

ERROR RESPONSE MODELS

A standardized error response model ensures every error looks the same, no matter which endpoint failed.

STANDARD ERROR RESPONSE — RFC 9457 PROBLEM DETAILS

```
{
  "type": "https://api.example.com/problems/validation",
  "title": "Validation failed",
  "status": 400,
  "detail": "One or more fields are invalid.",
  "instance": "/api/users",
  "code": "VALIDATION_FAILED",
  "correlationId": "3f9c2a6e7b4dd4e1a",
  "errors": [
    { "field": "email", "code": "EMAIL_INVALID",
      "message": "Enter a valid email address." }
  ]
}
```


KEY CONCEPTS

- Correlation ID vs. trace ID: correlationId is an application-level ID that links related logs/requests. A trace ID is a distributed-tracing ID for one trace. They may match in a simple app, but don't assume they're identical — observability platforms may add their own trace and span IDs.
- Timestamp is optional per your organization's standard, not a mandatory field.
- Design principles (from user-friendly error design): be clear and specific; include actionable detail (what field, what rule); link documentation by error type; keep tone professional and consistent across the API.

KEY TAKEAWAY A consistent error model — timestamp, status, error code, message, path, trace ID — makes APIs predictable and debuggable.

ERROR RESPONSE MODELS

1. SIMPLE ERROR MODEL




JSON RESPONSE EXAMPLE

```
{
  "timestamp": "2026-07-17T10:15:30Z",
  "status": 400,
  "error": "Bad Request",
  "message": "Invalid input provided"
}
```

USE CASE: Simple APIs, basic error information

2. DETAILED ERROR MODEL




JSON RESPONSE EXAMPLE

```
{
  "timestamp": "2026-07-17T10:15:30Z",
  "status": 422,
  "error": "Unprocessable Entity",
  "message": "Validation failed",
  "path": "/api/users",
  "errorCode": "VALIDATION_ERROR",
  "details": [
    { "field": "email", "message": "Email is invalid" },
    { "field": "age", "message": "Age must be >= 18" }
  ]
}
```

USE CASE: Validation errors, detailed field-level info

3. PROBLEM DETAILS (RFC 7807) MODEL




JSON RESPONSE EXAMPLE

```
{
  "type": "https://example.com/probs/invalid-input",
  "title": "Invalid Input",
  "status": 400,
  "detail": "The request contains invalid input",
  "instance": "/api/users/123",
  "extensions": {
    "errorCode": "INVALID_INPUT",
    "timestamp": "2026-07-17T10:15:30Z",
    "path": "/api/users/123"
  }
}
```

USE CASE: Standardized error format (RFC 7807)

4. APPLICATION SPECIFIC ERROR MODEL




JSON RESPONSE EXAMPLE

```
{
  "success": false,
  "data": null,
  "error": {
    "code": "USR_1001",
    "message": "User not found",
    "description": "No user exists with the given ID",
    "timestamp": "2026-07-17T10:15:30Z",
    "requestId": "a8f3d2e1-7b4c-4d2a-bc9f-123456789abc"
  }
}
```

USE CASE: Domain-driven applications, consistent envelope

5. ERROR MODEL WITH TRACE & SUPPORT INFO



JSON RESPONSE EXAMPLE

```
{
  "timestamp": "2026-07-17T10:15:30Z",
  "status": 500,
  "error": "Internal Server Error",
  "message": "An unexpected error occurred",
  "path": "/api/orders",
  "traceId": "4f3a1b2c9d4e7f8a",
  "support": {
    "contact": "support@example.com",
    "documentation": "https://docs.example.com/errors/500"
  }
}
```

USE CASE: Production APIs, easier debugging & support

HTTP STATUS CODES

Map exceptions to the right HTTP status code — it's part of your API contract.

CODE	MEANING	WHEN TO USE
400	Bad Request	Validation failures, malformed input
401	Unauthorized	Missing or invalid authentication
403	Forbidden	Authenticated but not authorized
404	Not Found	Resource doesn't exist
409	Conflict	Duplicate resource, state conflict
500	Internal Server Error	Unexpected server-side failure
503	Service Unavailable	Downstream dependency unavailable
422	Unprocessable Entity	Semantically invalid request under a standard that distinguishes it from 400
429	Too Many Requests	Client exceeded rate limit
502	Bad Gateway	Invalid response from downstream service
504	Gateway Timeout	Downstream timeout while acting as gateway/proxy

Follow the organization's API standard consistently. 400 is common; some APIs use 422 for semantic validation failures.

KEY TAKEAWAY Choosing the correct HTTP status code helps consumers handle errors programmatically and correctly.

HTTP STATUS CODES

CLASS	RANGE	DESCRIPTION	COMMON CODES		MEANING / WHEN TO USE
1xx Informational 	100 – 199	Request received, continuing process	100 Continue 101 Switching Protocols 102 Processing		Request received, client should continue Switching to a different protocol Processing request, no response yet
2xx Success 	200 – 299	Request successfully received, understood, and accepted	200 OK 201 Created 204 No Content 206 Partial Content		Request successful Resource created successfully Request successful, no content to return Partial content returned
3xx Redirection 	300 – 399	Further action needs to be taken by the client	301 Moved Permanently 302 Found 304 Not Modified 307 Temporary Redirect 308 Permanent Redirect		Resource permanently moved Resource temporarily moved Resource not modified, use cached version Temporary redirect, repeat request with new URL Permanent redirect, repeat request with new URL
4xx Client Error 	400 – 499	Request contains bad syntax or cannot be fulfilled	400 Bad Request 401 Unauthorized 403 Forbidden 404 Not Found 409 Conflict 422 Unprocessable Entity		Invalid request syntax Authentication required or failed Client doesn't have permission Requested resource not found Request conflicts with current state Request well-formed but cannot be processed
5xx Server Error 	500 – 599	Server failed to fulfill a valid request	500 Internal Server Error 501 Not Implemented 502 Bad Gateway 503 Service Unavailable 504 Gateway Timeout		Unexpected server error Server does not support the functionality Invalid response from upstream server Server is currently unavailable Gateway did not receive timely response

BUSINESS EXCEPTIONS

Business exceptions represent violations of domain rules — invalid operations, entity state issues, duplicates.

EXAMPLE

```
public abstract class BusinessException
    extends RuntimeException {

    private final String code;
    private final int status;
}
```

STATUS MAPPING BY MEANING

MEANING	TYPICAL STATUS
Duplicate resource	409
Invalid state transition	Often 409
Semantically invalid command	422 or 400
Insufficient funds	Domain-specific — often 409 or 422
Business authorization failure	May be 403

BEST PRACTICES

- Map status by meaning, not exception class name (see table below).
- Use specific exception types only where callers need distinct handling.
- A shared typed exception with a stable error code can suit related failures.
- Never expose internal state, stack traces, or raw exception messages to consumers.

USAGE EXAMPLE

```
throw new BusinessException(
    CustomerErrorCode.INVALID_STATUS_TRANSITION,
    "Customer cannot move from ACTIVE to PROSPECT");
```

Consumer-facing messages should not come directly from raw exception messages. The exception carries a stable internal code and safe parameters; the handler maps it to a consumer-safe message; logs retain the technical details; the response follows the documented error contract.

KEY TAKEAWAY Well-named business exceptions make your code self-documenting and your error responses meaningful.

BUSINESS EXCEPTIONS



COMMON BUSINESS EXAMPLES

SCENARIO	EXAMPLE
User related	User not found, user already exists, account locked
Order related	Insufficient stock, invalid order state, order cannot be canceled
Payment related	Payment declined, invalid payment method, limit exceeded
Resource related	Resource not available, quota exceeded, duplicate resource
Authorization related	Operation not allowed, invalid role for operation

BUSINESS EXCEPTION HIERARCHY

```
graph TD; A[BusinessException (Base Class)] --> B[UserException]; A --> C[OrderException]; A --> D[PaymentException]; B --> E[UserNotFound Exception]; C --> F[InsufficientStock Exception]; D --> G[PaymentDeclined Exception]; E --> H[InvalidState Exception]; F --> I[DuplicateResource Exception]; G --> J[QuotaExceeded Exception];
```

BEST PRACTICES

- Use meaningful exception names that reflect business rules
- Include clear, user-friendly messages
- Include error codes for consistent handling
- Do not expose internal details
- Log exception with full context for troubleshooting

EXAMPLE BUSINESS EXCEPTION RESPONSE

```
{
  "timestamp": "2026-07-17T10:15:30Z",
  "status": 400,
  "error": "Business Rule Violation",
  "message": "Insufficient stock for product with id 12345",
  "errorCode": "ORDER_STOCK_INSUFFICIENT",
  "path": "/api/orders",
  "details": { "productId": "12345", "requestedQty": 10, "availableQty": 2 }
}
```

FIELD	DESCRIPTION
timestamp	Time when the error occurred
status	HTTP status code
error	Short error summary
message	Human-readable error message
errorCode	Application-specific error code
path	API endpoint where error occurred
details	Additional context for the error

VALIDATION EXCEPTIONS

Module 14 covered Bean Validation in depth — here we map validation exceptions to a stable field-error response.

EXAMPLE HANDLER

```
@ExceptionHandler(MethodArgumentNotValidException.class)
public ResponseEntity<ErrorResponse> handleValidation(
    MethodArgumentNotValidException ex) {...}
```

EXAMPLE FIELD ERRORS (NESTED / INDEXED PATHS)

```
"errors": [
  { "field": "address.postalCode",
    "code": "FIELD_INVALID",
    "message": "Enter a valid postal code" },
  { "field": "items[2].quantity",
    "code": "FIELD_INVALID",
    "message": "Quantity must be positive" }
]
```

BEST PRACTICES


- Mapping pipeline: framework exception → field extraction → stable field-error model → status selection → correlation ID.
- Preserve a stable, consumer-understandable field path for nested and collection validation errors (see example, right) — not just a flat "field": "email".
- Follow the organization's API standard consistently. 400 is common; some APIs use 422 for semantic validation failures.

KEY TAKEAWAY Field-level validation errors help API consumers fix their requests quickly, without guesswork.

VALIDATION EXCEPTIONS



COMMON VALIDATION SCENARIOS	
SCENARIO	EXAMPLES
Required Field Missing	name, email, address
Invalid Data Type	Expected number but received string
Length/Size Constraint	Minimum/maximum length violated
Format Constraint	Invalid email, phone number, date format
Range Constraint	Value out of allowed range (min/max)
Uniqueness Constraint	Duplicate value for unique field



BEST PRACTICES	
	Validate input at the boundary (API layer)
	Return clear and specific validation messages
	Provide field-level error details
	Use consistent error codes for validation errors
	Do not expose internal validation rules

EXAMPLE VALIDATION ERROR RESPONSE

```
{
  "timestamp": "2026-07-17T10:15:30Z",
  "status": 400,
  "error": "Validation Failed",
  "message": "Request data validation failed",
  "path": "/api/users",
  "errorCode": "VALIDATION_ERROR",
  "validationErrors": [
    { "field": "email", "message": "Invalid email format", "rejectedValue": "john@com" },
    { "field": "age", "message": "Must be between 18 and 65", "rejectedValue": 15 },
    { "field": "name", "message": "Name is required", "rejectedValue": null }
  ]
}
```

FIELD	DESCRIPTION
timestamp	Time when the error occurred
status	HTTP status code (400 Bad Request)
error	Short error summary
message	Human-readable error message
path	API endpoint where error occurred
errorCode	Application-specific error code
validationErrors	List of field-level validation errors
field	Name of the field that failed validation
message	Validation error message for the field
rejectedValue	Value that was rejected

SECURITY EXCEPTIONS

Authentication and authorization failures need careful handling — informative, but not revealing.

EXCEPTION	MEANING	STATUS CODE
AuthenticationException	Missing or invalid credentials	401 Unauthorized
AccessDeniedException	Authenticated but lacks permission	403 Forbidden
TokenExpiredException	JWT/session token has expired	401 Unauthorized

BEST PRACTICES


- Never reveal whether a username or resource exists in error messages.
- Log security exceptions with context, but don't return internal details to the client.
- Return generic, consistent messages for authentication failures.
- Depending on the threat model, an API may intentionally return the same response for nonexistent and unauthorized resources (e.g., 404 instead of 403) — follow the security standard consistently.
- Expose only the minimum detail required by the client contract. OAuth-based APIs may use standardized authentication error fields while avoiding sensitive internals.

KEY TAKEAWAY Security exceptions must be handled carefully — informative to developers via logs, generic to the client.

SECURITY EXCEPTIONS



COMMON SECURITY EXCEPTIONS	
EXCEPTION	DESCRIPTION
AuthenticationException	Authentication failed (invalid credentials, token expired)
AuthorizationException	User not authorized to access resource
AccessDeniedException	Access is denied (insufficient permissions)
AccountLockedException	User account is locked
TokenExpiredException	Security token has expired
CSRFException	Invalid or missing CSRF token



BEST PRACTICES	
	Do not expose sensitive security details
	Use generic messages for security errors
	Log full details for monitoring and analysis
	Use proper HTTP status codes
	Prevent user enumeration in error messages
	Review and audit security logs regularly

EXAMPLE SECURITY ERROR RESPONSE

```
{
  "timestamp": "2026-07-17T10:15:30Z",
  "status": 401,
  "error": "Unauthorized",
  "message": "Authentication failed",
  "errorCode": "AUTHENTICATION_FAILED",
  "path": "/api/orders",
  "traceId": "4fa3a1b2c9d4e7f8a",
  "details": null
}
```

FIELD	DESCRIPTION
timestamp	Time when the error occurred
status	HTTP status code
error	Short error summary
message	Human-readable message
errorCode	Application-specific error code
path	API endpoint where error occurred
traceId	Unique ID for tracing the request
details	Additional details (if safe to expose)

LOGGING ERRORS

Good logging turns a production incident into a quick fix — log with context, not noise.



Correlation IDs

- Trace a request across services and logs.



Right Log Level

- Depends on the situation — see table below.



No Sensitive Data

- Never log passwords, tokens, or PII; use non-sensitive stable IDs.



Structured Logging

- JSON logs for easy parsing and search.



Include Context

- Sufficient diagnostic context, not everything available.



Monitor & Alert

- Set up alerts on error rate spikes.

SITUATION	TYPICAL LOGGING
Client validation error	No error stack trace; metric or debug if needed
Expected domain conflict	Info/debug or structured event
Security failure	Security audit event, carefully rate-limited
Downstream transient failure	Warn/error based on impact
Unexpected internal failure	Error with stack trace

USER IDS & DUPLICATE LOGGING

Use non-sensitive, stable identifiers where approved; avoid emails, tokens, personal data, and raw payloads. Hash or tokenize identifiers if policy requires.


Log once at the boundary that has enough context to act. Lower layers should enrich or rethrow, not repeatedly log the same stack trace.

KEY TAKEAWAY Log sufficient diagnostic context, not all available context — apply redaction, allowlists, and size limits so correlation IDs and structured logs stay fast and traceable.

LOGGING ERRORS



WHAT TO LOG	
!	Error type / exception class
📄	Error message
</>	Stack trace
👤	User ID / affected user
🔗	Request path & HTTP method
#	Correlation / Trace ID
📋	Request parameters (sanitized)
🕒	Timestamp
⚙️	Environment / Service name
💻	Host / IP address



BEST PRACTICES	
🎯	Log at the right level (ERROR, WARN, INFO)
📄	Use structured logging (JSON) for easy parsing
🚫	Never log sensitive data (passwords, tokens, PII)
🔗	Include correlation/trace ID for end-to-end tracking
📁	Centralize logs for analysis and retention
🕒	Set log retention and rotation policies
🔔	Monitor logs and set up alerts for critical errors

EXAMPLE STRUCTURED ERROR LOG (JSON)
<pre>{ "timestamp": "2026-07-17T10:15:30.123Z", "level": "ERROR", "message": "Failed to process user request", "errorType": "ValidationException", "stackTrace": "com.example.service.UserService.validate(UserService.java:45)\n...", "path": "/api/users", "method": "POST", "userId": "user123", "correlationId": "4fa3a1b2c9d4e7f8a", "requestId": "req-7890", "ip": "10.0.0.5", "service": "user-service", "environment": "production" }</pre>

LOG LEVELS		
LEVEL	DESCRIPTION	WHEN TO USE
ERROR	Serious issues that prevent request completion	Exceptions, failures, application errors
WARN	Potential issues that do not stop execution	Deprecated usage, retries, slow response
INFO	General operational events	Request received, process completed
DEBUG	Detailed diagnostic information	Development and troubleshooting
TRACE	Very detailed execution trace	Deep debugging, performance analysis

ENTERPRISE API ERROR STANDARDS

Consistent error standards across all services make your enterprise APIs predictable and easy to integrate.



Standard Error Schema

- Same fields across every API in the organization.



Error Code Registry

- Type-specific documentation links per code, not one generic URL.



Documented Errors

- Every error code documented in the API contract.



Traceability

- Correlation IDs propagated across services.



Security by Default

- No stack traces or internals in production responses.



Observability

- Metrics and dashboards for error rates by type.


ERROR-CODE REGISTRY GOVERNANCE

Define per code: ownership, naming convention, uniqueness, deprecation, documentation, mapping to HTTP status, retryability, consumer action, and security classification.


CODE	STATUS	RETRYABLE	CONSUMER ACTION
CUSTOMER_STATUS_CONFLICT	409	No	Refresh state and choose a valid transition

KEY TAKEAWAY Enterprise error standards make debugging faster and integration easier across every team and service.

ENTERPRISE API ERROR STANDARDS



ERROR CATEGORIES		
CATEGORY	DESCRIPTION	EXAMPLE
VALIDATION	Invalid input or business rule violations	Missing required field, invalid email format
AUTHENTICATION	Authentication failures	Invalid credentials, token expired
AUTHORIZATION	Insufficient permissions	Access denied to resource
BUSINESS	Business rule or process errors	Insufficient balance, order cannot be cancelled
SYSTEM	Server or dependency failures	Database unavailable, service timeout
EXTERNAL	Third-party or external service errors	Payment gateway error, upstream service failure



ERROR CODE DESIGN	
	Structured & Consistent Use a consistent format and naming convention.
	Meaningful Codes should be descriptive and easy to understand.
	Categorized Group codes by domain and error category.
	Unique Each error condition should have a unique code.
	Stable Once published, codes should remain stable.

```
EXAMPLE ERROR RESPONSE (JSON)

{
  "timestamp": "2026-07-17T10:15:30Z",
  "status": 400,
  "error": "Bad Request",
  "message": "Validation failed for one or more fields.",
  "errorCode": "VAL_1001",
  "path": "/api/users",
  "traceId": "4fa3a1b2c9d4e7f8a1234567890abcdef",
  "details": {
    "fieldErrors": [
      { "field": "email", "rejectedValue": "john@com", "message": "Invalid email format" },
      { "field": "age", "rejectedValue": -5, "message": "Age must be greater than 0" }
    ]
  }
}
```

BEST PRACTICES	
	Use appropriate HTTP status codes.
	Provide clear, user-friendly messages.
	Use consistent error response format.
	Maintain a catalog of error codes.
	Do not expose internal implementation details.
	Log errors with full context.
	Monitor and alert on error patterns.
	Review and update standards regularly.
	Document for developers and consumers.
	Ensure security and compliance.

LAB OVERVIEW – API EXCEPTION HANDLING

Design a consistent API error model for the Customer Management Platform, then map the pattern to Spring later.

WHAT YOU WILL BUILD

- ExceptionToErrorResponseMapper for validation and business failures (framework-neutral, not Spring's global interception).
- BusinessException factories for not-found and conflict cases.
- Standard ErrorResponse with status mapping and correlation IDs.
- Spring's @RestControllerAdvice delegates to the same mapping policy later.

TECH STACK

ITEM	VALUE
Language	Java 21
Framework	Java + Bean Validation
Logging	stdout / simple logger
Testing	JUnit 5 + Maven
Pattern	Two-stage: framework-neutral mapper now, Spring adapter preview later

HIGH-LEVEL FLOW



CONVENTION

Domain/service methods may throw typed exceptions; the facade catches and maps them and returns ApiResponse. A layer should not both return a failure value AND throw for the same expected condition.

KEY TAKEAWAY Expected outcome: every failure path returns the same safe JSON shape with the right status and correlation ID.

LAB TASKS AND DELIVERABLES

Implement the lab's plain-Java error model, facade mapping, demos, and tests.

#	TASK	DETAILS	FORMAT
1	ExceptionToErrorResponseMapper	Map validation/business/unexpected failures (framework-neutral mapper).	.java
2	BusinessException	Add not-found and conflict factory methods.	.java
3	Standard Error Model	Include status, message, errors, and correlationId.	.java
4	Facade Integration	Return ApiResult.Ok/Fail consistently — don't also throw for the same expected condition.	.java
5	CLI Error Demos	Demonstrate 400, 404, and 409 JSON safely.	.java
6	Tests and Docs	Add JUnit 5 handler tests and document mappings.	Tests / docs

SUCCESS CRITERIA

- All failure paths share one JSON shape; 400/404/409 mappings are tested; correlationId is present; no stack traces or entity details leak.
- 500 unexpected-failure test: generic safe message, correlationId present, no stack trace in the JSON response, stack trace retained in the internal log.
- Field validation tests: stable field path, stable error code, safe message, multiple violations, deterministic ordering (if required), no entity-serialization leakage.
- Correlation ID: accept an approved incoming header or generate one; validate length/characters; include it in logs and response; propagate downstream; never trust it for identity or authorization.
- Serialize errors with the project's JSON serializer — never hand-concatenate untrusted strings into JSON.

KEY TAKEAWAY A robust, secure, well-documented exception handling strategy makes your API easy to integrate and trust.

API ERROR HANDLING — DEFINITION OF DONE

Use this checklist to confirm an exception handling strategy is complete, consistent, and safe before calling it done.

- 1 Errors are classified by source and retryability.
- 2 Every public error has a stable code.
- 3 HTTP status matches the API standard.
- 4 Response shape follows one documented schema.
- 5 No stack trace or sensitive data is exposed.
- 6 Validation errors contain safe field paths.
- 7 Dependency failures are mapped consistently.
- 8 Correlation ID is propagated end-to-end.
- 9 Unexpected failures are logged once with diagnostics.
- 10 Expected client errors do not create noisy error logs.
- 11 Error codes are documented and tested.
- 12 Generic fallback behavior is defined.

KEY TAKEAWAY Reliable APIs handle failure gracefully — consistent, secure, and transparent error handling builds consumer trust.

KNOWLEDGE CHECK

Test your understanding of exception handling in distributed APIs.

1. A downstream inventory service times out while your API acts as a gateway. Likely response?

- A. 200 OK, and retry silently
- B. 504 Gateway Timeout (architecture/policy-dependent)**
- C. 401 Unauthorized
- D. 201 Created

3. What should an error response NEVER include?

- A. A timestamp
- B. A trace ID
- C. Internal stack traces**
- D. An error message

2. Which status code fits a validation failure?

- A. 200 OK
- B. 400 Bad Request**
- C. 404 Not Found
- D. 500 Internal Server Error

4. Which error should normally avoid an ERROR-level stack trace?

- A. An unexpected null pointer exception**
- B. A downstream service crash
- C. A normal client validation failure**
- D. A corrupted database record

KEY TAKEAWAY Map failures to status by meaning, centralize handling, and never expose stack traces to API consumers.