

MODULE 21

API Observability and Monitoring

Explore how to gain deep visibility into the health, performance, and behavior of APIs in production.





MODULE 21 OVERVIEW

Build, observe, diagnose, and improve — observability is the foundation of production-ready APIs.

- This module explores how to gain deep visibility into APIs in production using logs, metrics, and traces, plus Spring Boot Actuator, alerting, and dashboards.

DURATION & LAB



Theory

Three pillars, Actuator, tracing, alerting, dashboards.



Lab

Add Actuator health checks and Micrometer metrics.



Scenario

Instrument, visualize, alert, and analyze issues.

BUILD

Metrics, logs, and traces via Actuator + Micrometer

OBSERVE & DIAGNOSE

Dashboards, alerting, distributed tracing

THE PAYOFF

Faster detection, faster root cause, higher reliability

KEY TAKEAWAY Good observability leads to faster insights and better outcomes — you can't improve what you can't see.

WHY OBSERVABILITY MATTERS

In modern distributed systems, issues are inevitable — observability gives you deep visibility to find and fix them fast.



Faster Issue Detection



Root Cause Analysis



Improved Reliability



Better User Experience



Data-Driven Decisions




Business Impact

Impact chain: Operational Impact (efficient ops) + Engineering Impact (faster debugging) + Customer Impact (reliable APIs) roll up to Business Impact — lower costs, better reliability, greater customer trust.

KEY TAKEAWAY You can't improve what you can't see — observability turns invisible failures into actionable insights.

API OBSERVABILITY

Gain full visibility into API health, performance, and user experience



MONITORING VS OBSERVABILITY

Monitoring watches defined signals and conditions. Observability is a system property: how well internal behavior can be inferred from emitted telemetry.

MONITORING

- Practice of watching defined signals, dashboards, and conditions.
- Key question: "Are our defined service objectives and known conditions healthy?"
- Predefined metrics, thresholds, and alerts.
- One of the ways teams USE observability data.

VS

OBSERVABILITY

- System property: how well internal behavior can be inferred from telemetry.
- Key question: "Do we have enough telemetry and context to investigate unexpected behavior?"
- Dynamic exploration using logs, metrics, and traces together.
- Instrumentation → telemetry → monitoring, alerting, investigation, capacity planning.

Monitoring is one way teams use observability data — not a competing approach. Instrumentation produces telemetry; telemetry feeds monitoring, alerting, investigation, and capacity planning.

KEY TAKEAWAY Monitoring watches defined signals for known conditions. Observability is whether your telemetry can answer the unknown questions.

API Monitoring Architecture

Real-time Visibility. Proactive Alerts. Reliable APIs.



THREE PILLARS OF OBSERVABILITY

Logs, metrics, and traces are the three core telemetry signals taught in this module. Mature observability platforms may also add profiles, events, and user-experience telemetry.

LOGS

What: Discrete events that happen in your system.

Answers: What happened?

Best for: Debugging, auditing, troubleshooting.

METRICS

What: Aggregated numerical data on health and performance.

Answers: How is the system performing?

Best for: Alerting, dashboards, capacity planning.

TRACES

What: End-to-end request flows across services.

Answers: Where is the issue happening?

Best for: Root cause analysis, latency tracking.

KEY TAKEAWAY Logs, metrics, and traces provide complementary evidence — consistent instrumentation and operational practice turn them into insight.

The Three Pillars of Observability

Observability is built on 3 types of telemetry data that help you understand your system.



LOGS

What happened?

Discrete events recorded at a specific point in time.
Best for understanding detailed behavior.

Examples

- Error messages
- Exception stack traces
- User activities
- Audit events

```
{
  "timestamp": "2026-07-17T10:15:30Z",
  "level": "ERROR",
  "service": "order-service",
  "message": "Payment failed",
  "orderId": "12345",
  "userId": "u789"
}
```

Best used for

- ✓ Investigating issues
- ✓ Root cause analysis
- ✓ Auditing and compliance



Analogy

Logs are like detailed news articles about what happened.



METRICS

What is happening?

Aggregated numerical measurements over time.
Best for understanding trends and system health.

Examples

- Request rate (RPS)
- Error rate (%)
- Latency (p50, p95, p99)
- CPU / Memory usage
- Throughput



Best used for

- ✓ Monitoring system health
- ✓ Detecting anomalies
- ✓ Capacity planning & alerting



Analogy

Metrics are like dashboards showing the speed and health of your car.



TRACES

Why is it happening?

Requests flow through your system and across services.
Best for understanding dependencies and latency.

Examples

- End-to-end request
- Spans across services
- Latency breakdown
- Dependencies



Best used for

- ✓ Understanding request flow
- ✓ Finding performance bottlenecks
- ✓ Analyzing service dependencies



Analogy

Traces are like GPS showing the full journey and where delays occur.

Stronger Together

 Logs tell you what happened.

+

 Metrics show you the impact and trends.

+

 Traces explain why it happened and where.

=

 Full observability gives you the complete picture to understand and improve your system.



Observability is not about more data, it's about the right data, correlated to answer the right questions.

LOGS

Module 20 covered structured logs, correlation IDs, privacy, and centralized logging in depth. Here, logs are correlated with metrics and traces during diagnosis.

EXAMPLE LOG ENTRY (JSON)

```
{
  "timestamp": "2025-05-20T10:15:24.123Z",
  "level": "INFO", "service": "order-service",
  "traceId": "a1b2c3d4e5f67890", "spanId": "b1c2d3e4f5678901",
  "message": "Order created successfully", "userId": "u12345",
  "orderId": "ord-98765", "status": "SUCCESS", "durationMs": 142
}
```

CROSS-SIGNAL WORKFLOW



- Workflow: Metric alert -> trace ID -> related structured logs.
- Include only approved, non-sensitive identifiers in logs. Do not log names, emails, tokens, session secrets, or regulated data.

Log levels (low to high severity): TRACE -> DEBUG -> INFO -> WARN -> ERROR. Typical flow: Application -> Log Collector (Filebeat/Fluentd) -> Log Store (Elasticsearch/CloudWatch) -> Visualization (Kibana/Grafana).

KEY TAKEAWAY Logs still matter, but in this module they are correlated with metrics and traces, not learned again from scratch.

LOGS

Record discrete events that help you understand what happened in your system.



METRICS

Metrics are numerical measurements collected over time, giving quantifiable insight into health, performance, and behavior.

METRIC	DESCRIPTION	TYPE	EXAMPLE
HTTP Request Rate	Number of requests per second	Counter	125 req/s
Response Time (p95)	95th percentile of API response time	Derived (timer/histogram)	480 ms
Error Rate	Percentage of failed requests	Derived (counters)	1.23 %
CPU / Memory Usage	Resource utilization	Gauge	62% / 1.4 GB

RED (RATE, ERRORS, DURATION) + USE (UTILIZATION, SATURATION, ERRORS) — DISTINCT FROM GOLDEN SIGNALS

**RED Method**
Rate, Errors, Duration — for request-driven services.

**USE Method**
Utilization, Saturation, Errors — for infrastructure/resources.

**Drive Alerting**
Use metric thresholds to trigger alerts and automation.

Metric types: Counter (increases only; resets on restart — query rates over time, not raw restart-to-restart values), Gauge (up/down), Timer/Histogram (duration distribution; preferred for cross-instance percentiles), Summary (client quantiles — may not aggregate across replicas).



Business examples: creation success/rejection-by-reason, lookup success/not-found, processing duration, readiness state. Avoid as labels: customer ID, email, request ID, trace ID, raw error text.

Low-cardinality labels: result=success, operation=create, reason=duplicate, route=/customers/{id}. Dangerous: customerId, traceId, errorMessage, full URI — risk: excess series, memory, cost, instability.

KEY TAKEAWAY Instrument counters, timers, and gauges correctly, keep labels low-cardinality, and combine metrics with logs and traces.


METRICS

Quantitative measurements over time that help you understand what is happening.



DISTRIBUTED TRACING FUNDAMENTALS

A trace shows how one request flows through multiple services. Spans capture timing and metadata, and context propagates across service boundaries.



A trace is made up of spans, each with a start time, end time, duration, and metadata. End-to-end latency follows the critical path — child-span durations may overlap, so do not always add every span duration together.

TRACE, SPAN & CONTEXT PROPAGATION

- Trace = the end-to-end journey of a request. Span = one unit of work within it; child spans reference a parent span ID (parent-child hierarchy).
- Continue an incoming W3C trace context when present; otherwise create a new trace at the system boundary — not at every entry point.
- Do not attach request bodies, tokens, raw SQL with personal values, or full customer records to spans — use route template, HTTP status, operation name, and safe error type instead.

Roles: OpenTelemetry provides instrumentation and context; the OpenTelemetry Collector handles collection/export; Jaeger, Zipkin, and vendor APM platforms serve as backends. Sampling: head sampling, tail sampling, and error-biased retention balance cost — not every request has a stored trace.


KEY TAKEAWAY Traces reveal latency and root cause across services — protect sensitive span data and sample wisely.

TRACES

Follow a request across services. Understand every step. Fix faster.

WHAT ARE TRACES?

Traces show the path of a request as it travels through multiple services. Each trace is made up of spans that represent individual operations.



KEY CONCEPTS

- Trace ID** – Unique identifier for a request across services
- Span** – A unit of work within a trace (e.g., database call)
- Parent Span** – The operation that triggered another
- Attributes** – Key-value data about the operation
- Duration** – Time taken by a span

HOW TRACES WORK




Each service creates a span and adds context so the entire request can be stitched together.

BENEFITS

- ✓ End-to-end visibility across services
- ✓ Identify slowdowns and bottlenecks
- ✓ Debug errors across distributed systems
- ✓ Understand dependencies and flows
- ✓ Improve performance and reliability


DISTRIBUTED TRACE EXAMPLE



SPAN LEGEND

- Service A 120 ms
- Service B 250 ms
- Service C 200 ms
- Service D 150 ms
- Parent → Child

TRACE IN ACTION




SPAN DETAILS (EXAMPLE)

Span ID	Operation	Service	Start Time	Duration	Status	Attributes (Key→Value)
a1b2c3d	GET /orders/123	Order Service	10:00:00.100	120 ms	✓	http.method=GET http.url=/orders/123
d4e5f6g	GET /inventory/123	Inventory Service	10:00:00.220	250 ms	✓	db.system=postgres rows=1
h7i8j9k	POST /payments	Payment Service	10:00:00.470	200 ms	✓	provider=stripe status=success
l1m2n3o	POST /email/send	Email Service	10:00:00.670	150 ms	✓	provider=smtp to=user@example.com

COMMON SPAN ATTRIBUTES

- http.method – Request method
- http.url – Request URL
- status.code – HTTP status code
- db.system – Database type
- peer.service – Calling service
- error – Error details (if any)
- user.id – User identifier
- env – Environment (prod, staging)

TOOLS & INTEGRATIONS



BEST PRACTICES

- ✓ Instrument all critical services
- ✓ Propagate trace context (W3C Trace Context)
- ✓ Sample intelligently (adaptive sampling)
- ✓ Capture meaningful attributes
- ✓ Monitor and alert on slow traces and errors
- ✓ Review traces regularly for optimization

SPRING BOOT ACTUATOR OVERVIEW

Actuator provides low-effort operational endpoints and baseline instrumentation. Production monitoring still requires configuration, security, and operational design.

ENDPOINT	PURPOSE
/actuator/health	Health status of the application
/actuator/info	Application information
/actuator/metrics	Application metrics
/actuator/prometheus	Metrics in Prometheus format
/actuator/env	Environment properties

ENABLING ACTUATOR (application.yml)

```
<dependency><artifactId>spring-boot-starter-actuator</artifactId></dependency>
<dependency><groupId>io.micrometer</groupId><artifactId>micrometer-registry-prometheus</artifactId></dependency>
management:
  endpoints: {web: {exposure: {include: health,info,metrics,prometheus}}}
  endpoint: {health: {probes: {enabled: true}, show-details: when_authorized}}
```

RESTRICTED DIAGNOSTIC ENDPOINTS — /actuator/env, /beans, /mappings, /loggers, /threaddump: never expose publicly; require authN/authZ, least privilege, and audited access (a separate management port/network is ideal). exposure.include controls WEB EXPOSURE; security config decides WHO can access.

KEY TAKEAWAY Actuator is a low-effort start, not a finished solution — secure, expose deliberately, and add real instrumentation.

Spring Boot Actuator

Production-ready features to monitor and manage your Spring Boot applications.



Spring Boot Actuator provides built-in endpoints to monitor application health, metrics, info, and more. It helps you operate, monitor, and manage your application in production.

1. ENABLE ACTUATOR

Maven

```
<dependency>
  <groupId>org.springframework.boot</groupId>
  <artifactId>spring-boot-starter-actuator</artifactId>
</dependency>
```

Gradle

```
implementation
  'org.springframework.boot:spring-boot-
  starter-actuator'
```

application.properties / application.yml

```
management.endpoints.web.exposure.include=*
management.endpoint.health.show-details=when_authorized
```

2. ACTUATOR ARCHITECTURE

Your Spring Boot Application



Actuator Endpoints



Output / Consumers



Actuator Core



Endpoint Discovery



Security



Endpoint Web Extension



Configuration Properties



Security Note: Actuator endpoints can expose sensitive information. Always secure them in production using authentication and authorization.

Base Path (default)

/actuator

Example:

<http://localhost:8080/actuator/health>

3. COMMON ACTUATOR ENDPOINTS

Endpoint	Path	Description
Health	/actuator/health	Application health status
Metrics	/actuator/metrics	Application metrics (counters, gauges, etc.)
Info	/actuator/info	Application information
Beans	/actuator/beans	List of Spring beans
Mappings	/actuator/mappings	Request mapping details
Environment	/actuator/env	Environment properties
Config Props	/actuator/configprops	Configuration properties
Loggers	/actuator/loggers	View and modify logging levels
Threads	/actuator/threads	Thread dump
Auditevents	/actuator/auditevents	Security audit events
Shutdown*	/actuator/shutdown	Graceful shutdown (disabled by default)

4. HOW IT WORKS



1. Request
Client requests an actuator endpoint



2. Endpoint Handler
Actuator maps the request to the appropriate endpoint



3. Data Collection
Gathers data from application context, metrics registry, etc.



4. Response
Returns data in JSON (or other supported format)

5. CUSTOMIZATION

- Include/Exclude Endpoints

```
management.endpoints.web.exposure.include=health,info,metrics
```
- Change Base Path

```
management.endpoints.web.base-path=/manage
```
- Customize Health Indicators
Create custom HealthIndicator beans
- Tagging & Additional Info
Expose custom info and metrics

6. BENEFITS

- ✓ Built-in monitoring and management
- ✓ Production-ready insights
- ✓ Seamless integration with monitoring tools
- ✓ Easy to extend and customize
- ✓ Helps in faster troubleshooting
- ✓ Improves application reliability

INTEGRATION ECOSYSTEM



* The /actuator/shutdown endpoint is disabled by default and should only be enabled in controlled environments.

HEALTH ENDPOINTS

Health endpoints provide real-time status about your application and its dependencies — liveness and readiness answer different questions.

ENDPOINT	DESCRIPTION	TYPICAL USE
/actuator/health	Overall health status	Basic health check
/actuator/health/liveness	Is the application alive?	Kubernetes liveness probe
/actuator/health/readiness	Is it ready to serve traffic?	Kubernetes readiness probe
/actuator/health/db	Database health	Check DB connectivity

EXAMPLE RESPONSE (AUTHORIZED VIEW)

```
{ "status": "UP",  
  "components": { "db": { "status": "UP" },  
                  "diskSpace": { "status": "UP" } },  
  "groups": ["liveness","readiness"] }
```

BEST PRACTICE

- Liveness: is the process stuck/irrecoverable and should it restart? Avoid checking external dependencies that may recover on their own.
- Readiness: can this instance serve traffic now? A DB outage should usually fail readiness, not liveness.

Health probes signal orchestrators and load balancers to route or recover traffic — the endpoint itself does not perform recovery. Component subpaths (e.g. /health/db) vary by Spring Boot version; verify for yours.
Custom readiness checks must stay fast and bounded: no expensive remote calls, long database queries, or full business transactions on every probe.

KEY TAKEAWAY Liveness and readiness answer different questions — health signals drive orchestrator decisions, they don't make them.

METRICS COLLECTION

Micrometer is a metrics and observation instrumentation facade that Spring Boot uses — much more than just an export library.



MICROMETER REGISTRIES (EXPORTS) — PROMETHEUS IS ONE OPTION



Prometheus

Pull-based: Prometheus scrapes /actuator/prometheus on an interval.



Datadog / New Relic

Push-based: the application/agent exports telemetry to the vendor.



CloudWatch

Push-based, AWS-native monitoring and export.


```
Timer.Sample s = Timer.start(registry); try { return service.createCustomer(req); } finally { s.stop(Timer.builder("crm.customer.operation").tag("operation","create").register(registry)); }
```

Use finally so duration records even on failure. Prefer pre-registering meters in a dedicated component (e.g. CustomerMetrics) rather than scattering builders through business-logic code.

KEY TAKEAWAY Micrometer is Spring Boot's instrumentation facade — Prometheus is one export option, not the only one.

SERVICE PERFORMANCE SIGNALS

Track request-level signals (RED) plus resource saturation using percentiles, not averages, to catch degradation before users notice.



PERCENTILES, SLOs & DIAGNOSIS


- Prefer percentiles (p50, p95, p99) over averages — watch request volume alongside latency; averages hide outliers.
- SLI = measured indicator (e.g. successful-request %); SLO = target (e.g. 99.9% over 30 days); SLA = contractual commitment; error budget = allowed unreliability.
- Diagnosis example: p95 latency rises while error rate is normal and CPU saturation climbs — investigate the saturated dependency before changing code.
- Correlate metrics with logs and traces to confirm root cause quickly.

Example alert: burn rate shows the error budget is being consumed too fast — page before the SLO is breached, not after. Common tools: Prometheus, Grafana, Actuator, Micrometer, Datadog, CloudWatch.

KEY TAKEAWAY Measure with percentiles, define SLOs, and correlate signals — performance monitoring is proactive, not reactive.

ALERTING FUNDAMENTALS

Alerting helps you proactively detect and respond to issues, notifying the right people at the right time.



ALERT SEVERITY LEVELS

LEVEL	EXPECTED ACTION	EXAMPLE
Page	Immediate human response	High error rate, service unavailable
Ticket	Investigation during working hours	High latency, high CPU usage
Dashboard/event	No immediate action (often not an alert at all)	Deployment successful, scale up

Severity (page/ticket/dashboard) and state (firing, acknowledged, resolved) are different dimensions — resolved is a state, not a severity level.

Avoid ambiguous "MTTR" — use explicit terms: mean time to detect, mean time to acknowledge, mean time to restore service, mean time to resolve root cause.

Every actionable alert needs an owner/team, runbook, dashboard link, query, and escalation route — not just a threshold. Notification channels: Slack, Email, PagerDuty, Teams, Opsgenie, Webhook, SMS.

KEY TAKEAWAY Good alerts map severity to action, separate state from severity, and always have an owner and a runbook.

MONITORING DASHBOARDS

A dashboard should answer specific operational questions, not just display KPI tiles.

KPI TILE	EXAMPLE
Total Requests (RPS)	1,256
p95 Response Time	215 ms
Error Rate	0.35%
Current 30-Day Availability	99.95%

DASHBOARD TYPES & BEST PRACTICES

- A good dashboard answers: is the service healthy, are users impacted, what changed, where to investigate, and is capacity constrained?
- Organize by audience: executive/business, service overview, on-call diagnosis, capacity planning, release comparison.
- Show SLO target and error-budget remaining next to current availability — don't blend them into one 'SLA' number.
- For accessibility, mark thresholds with labels, icons, or patterns, not color alone.

Dashboards also track Active Users (e.g. 3,482). Common tools: Grafana, Kibana, Prometheus, Datadog, New Relic, CloudWatch.

KEY TAKEAWAY A good dashboard answers real questions for its audience — accessible status, not just a KPI wall.

LAB OVERVIEW – OBSERVABILITY AND MONITORING

Add Spring Boot Actuator health checks and Micrometer metrics to the Northstar CRM.

WHAT YOU WILL LEARN

- Expose /actuator/health with liveness and readiness probes.
- Register Micrometer metrics (crm.customer.operations, crm.customer.operation.duration) for CRM create/get.
- Add a CrmReadinessIndicator (a labeled training simulation) that fails independent of liveness.
- Automate Actuator smoke tests and write a monitoring report.

LAB OUTCOME

- You will expose Actuator health/readiness/liveness probes, register low-cardinality Micrometer metrics for CRM create/get, and document production exposure restrictions. This is a production-ready foundation for health and metrics — not full tracing, a metrics backend, dashboards, or alert delivery.

LAB ENVIRONMENT

ITEM	VALUE
Application	Northstar CRM (single service)
Monitoring Stack	Actuator + Micrometer
Tools	curl, Maven, IntelliJ IDEA
Time	45 minutes

Prerequisites: Labs 19-20 complete (CRM create/get API, structured logging), JDK 21, Maven 3.9+, IntelliJ IDEA Community. Scope: this lab exposes /actuator/metrics only — Prometheus scraping is instructor-provided/demo, not built here.

KEY TAKEAWAY This lab builds the exact Actuator health and metrics foundation production monitoring depends on.

LAB TASKS AND DELIVERABLES

Add Actuator health probes and Micrometer metrics to the Northstar CRM.

#	TASK	KEY ACTIVITIES	FOCUS
1	Copy Lab 20 & Add Actuator	Copy lab20-crm to lab21-crm; add the Actuator starter.	Setup
2	Configure Health & Metrics	Expose health, info, and metrics via application.yml.	Configuration
3	Verify Liveness & Readiness	curl /actuator/health/liveness and /health/readiness.	Probes
4	Add CrmReadinessIndicator	Indicator flips readiness OUT_OF_SERVICE; liveness UP.	Readiness
5	Register Create/Get Metrics	Wire CustomerMetrics counters/timers by result tag.	Metrics
6	Drive Traffic & Metrics	POST/GET CUS-1001; compare metric values within the same running instance.	Verification
7	Automate Actuator Tests	ActuatorIT verifies liveness UP, readiness OUT_OF_SERVICE, and tagged metric increments.	Testing
8	Write Monitoring Report	Document metrics, alert idea, and exposure limits.	Documentation

SUCCESS CRITERIA

- Readiness can go OUT_OF_SERVICE (a labeled training simulation) while liveness stays UP; create-success counters increase after POST; ActuatorIT passes; monitoring report documents metrics, alert idea, and production exposure restrictions.
- Metrics: crm.customer.operations, crm.customer.operation.duration (tags operation=create|get, result=success|not_found|validation_error|failure — never customerId/correlationId/raw error text). ActuatorIT also checks tag-only increments and restricted exposure; timers record duration via finally even on failure; avoid a redundant duplicate counter.
- Production exposure: health limited, probes cluster-internal, metrics operator-only, env/beans/loggers restricted (scope: /actuator/metrics only — Prometheus is instructor-provided/demo). Report includes a programmatic tag-cardinality check and one actionable alert with signal, threshold, window, severity, owner, and runbook.

KEY TAKEAWAY You will have Actuator health probes and low-cardinality Micrometer metrics an operator can trust — with an explicit, honest scope.

PRODUCTION MONITORING: DEFINITION OF DONE

Production monitoring for an API is "done" when every item below holds true — prioritized by user impact, service objectives, critical dependencies, saturation, business workflows, then security/compliance.

- 1 User-facing SLIs are defined.
- 2 SLO targets and error-budget windows are documented.
- 3 Metrics use controlled, low-cardinality labels.
- 4 Liveness and readiness have distinct semantics.
- 5 Actuator exposure is restricted and authenticated.
- 6 Dashboards answer real operational questions.
- 7 Alerts are actionable, owned, and mapped to severity.
- 8 Runbooks exist for every actionable alert.
- 9 Logs, metrics, and traces share context (trace/correlation IDs).
- 10 Telemetry cost, alert tests, and post-incident reviews happen regularly.

KEY TAKEAWAY Monitoring is 'done' when this checklist holds true, not when a dashboard merely exists.

KNOWLEDGE CHECK (1 OF 2)

Test your understanding of observability and monitoring concepts.

1. A team adds a metric label `customerId=CUS-1001`. Why is this risky?

- A. No effect on the metrics system
- B. Creates high-cardinality series, exposes IDs**
- C. Makes dashboards load faster
- D. Required for Prometheus to work

3. Which library provides Spring Boot's metrics instrumentation facade?

- A. Logback
- B. Micrometer**
- C. Zipkin
- D. Grafana

2. What is the primary purpose of distributed tracing?

- A. Store application logs
- B. Track a request across multiple services**
- C. Visualize server hardware metrics
- D. Trigger alerts for log errors

4. The database becomes temporarily unavailable. What should usually happen to the liveness probe?

- A. Immediately fail and restart the pod
- B. Stay UP; readiness should fail instead**
- C. No relationship to database health
- D. Return HTTP 500 to callers

KEY TAKEAWAY Metrics, logs, and traces are the three core signals — Actuator and Micrometer are how you collect them.

KNOWLEDGE CHECK (2 OF 2)

Four more questions, plus the full answer key.

5. Which is an example of a good alert rule?

- A. CPU usage > 50%
- B. Error rate > 5% for 5 minutes**
- C. Memory used > 10%
- D. Any 4xx error occurs

7. How do monitoring and observability differ?

- A. They are the same thing
- B. Monitoring = what; Observability = why**
- C. Observability only works in development
- D. Monitoring replaces the need for logs

ANSWER KEY

- 1-B 2-B 3-B 4-B 5-B 6-C 7-B 8-B

6. Which component visualizes metrics and dashboards?

- A. Alertmanager
- B. Prometheus
- C. Grafana**
- D. Spring Cloud Sleuth

8. Why secure Actuator endpoints in production?

- A. They are never needed
- B. They expose sensitive operational information**
- C. They slow down the application
- D. They only work with Prometheus

KEY TAKEAWAY Observability is not a one-time setup — continuously measure, learn, and improve your systems.

WEEK 2 REVIEW AND KEY TAKEAWAYS

This week, you built a strong foundation in enterprise Java development — organized into five connected themes, not a module-by-module list.

THIS WEEK'S THEMES

- ✓ **Build & dependency management (Modules 8-9)**
- ✓ **AI-assisted engineering (Modules 10-11)**
- ✓ **API & service-layer design (Modules 12-16)**
- ✓ **Testing strategy (Modules 17-19)**
- ✓ **Logging & observability (Modules 20-21)**

KEY TAKEAWAYS

- Build right from the start with a well-structured project and layered architecture.
- Design robust APIs with clear contracts and validated inputs.
- Test at every level: unit, integration, and UI.
- Observe and act — use logging, metrics, and tracing to detect issues early.
- Automate builds and use AI tools like GitHub Copilot responsibly.

WHAT'S NEXT? — Week 3 dives into Spring Framework, Data Access with JPA, Cloud-Native Development, and CI/CD Deployment.

KEY TAKEAWAY You are now equipped to build, test, and operate robust, scalable, and maintainable enterprise systems.

"You can't improve what you can't see. Good observability leads to faster insights and better outcomes."

MODULE 21 COMPLETE • WEEK 2 COMPLETE

API Observability and Monitoring

You can now build a practical, production-ready observability strategy using industry-standard tools and practices.